



MiPhi Semiconductors Pvt Ltd

GUI Instruction manual

MiPhi Semiconductors Pvt Ltd

Tel: +91 89044 05190 // +91 80421 66207

E-mail: inquire@miphi.in

ALL RIGHTS ARE STRICTLY RESERVED. ANY PORTION OF THIS PAPER SHALL NOT BE REPRODUCED, COPIED, OR TRANSLATED TO ANY OTHER FORMS WITHOUT PERMISSION FROM MIPHI SEMICONDUCTORS PVT LTD.

MiPhi may make changes to specifications and product description at any time without notice. MIPHI and the MiPhi logo are trademarks of MiPhi Semiconductors Pvt Ltd, registered in the United States and other countries. Products and specifications discussed herein are for reference purposes only. Copies of documents which include information of part number or ordering number, or other materials may be obtained by emailing us at inquire@miphi.in.

©2025 MiPhi Semiconductors Pvt Ltd. All Rights Reserved

CONFIDENTIAL

Table of Contents

1. Environment Preparation.....	4
1.1 Supported OS	
1.2 Install and start Docker	
1.3 NVIDIA Container Toolkit	
2. Retrieve Docker .tar Files.....	4
3. Load Docker Images.....	5
3.1 Prepare & Import Docker Image Archives	
3.2 Verify the Docker Images	
3.3 Check Docker Network	
4. Run GUI Container.....	6
4.1 Parameters to Update	
4.2 Example Parameter Values	
4.3 Verify Container Status	
5. Run GPU Utility Container.....	10
5.1 Parameters to Update	
5.2 Example Parameter Values	
5.3 Verify GPU Container	
6. Post-Deployment Checklist.....	12
6.1 Verify running containers	
6.2 Access the GUI	
6.3 Confirm GPU access	

1. Environment Preparation

1.1 Supported OS

- Ubuntu 22.04.4 LTS - using this version helps avoid unexpected compatibility issues.

1.2 Install and start Docker

- Docker bundles your front end, back end, and database into one package.
- You start the entire setup with a single command i.e. **docker run**, so you don't need to install each part separately or worry about version mismatches.

1.3 NVIDIA Container Toolkit

- NVIDIA Container Toolkit installed for GPU support inside Docker
- This makes heavy tasks run much faster than if they only used the regular processor.

2. Retrieve Docker .tar Files

Download the Docker image .tar files directly into the server's working directory via wget, ensuring they're ready for immediate loading

```
wget http://<server>/stream.tar  
wget http://<server>/gpu.tar
```

3. Load Docker Images

3.1 Prepare & Import Docker Image Archives

Make sure you're in the **same directory** where you ran your wget commands (where stream.tar and gpu.tar were downloaded).

- Import the GUI image:
 - The stream.tar file is a complete snapshot of your GUI application—including the web interface, backend services, and database schema—all pre-configured and packaged together.
 - Loading it into Docker unpacks that package, calls it **integration-stream**, and lets you start everything with a single command.

```
docker load -i stream.tar
```

- Import the GPU image:
 - The gpu.tar file provides the tools your container needs to read GPU metrics—like DRAM and memory usage—so you can keep an eye on performance.
 - Loading it makes the **testing-gpu** image available in Docker, ready for use by your GPU utility container.

```
docker load -i gpu.tar
```

3.2 Verify the Docker Images

- Run the following command to list Docker images:
 - Make sure both your GUI and GPU images are ready to use before starting any containers.

```
docker images
```

You should see:

- integration-stream (GUI image)

```
Last login: Mon May 12 14:35:32 2025 from 192.168.100.1
miphi@miphi:~$ docker images
REPOSITORY              TAG               IMAGE ID           CREATED            SIZE
nvme-inference           latest            37e5e06a3060      4 days ago        28.9GB
gui-without-supervisor   latest            c2f89c4ca7a6      6 days ago        24GB
milestone-10-dev         latest            ea032e4f159f      9 days ago        23.9GB
integration-stream        latest            32ce225f1750      9 days ago        24GB
milestone-9              latest            a80b5ed15f17      9 days ago        23.9GB
nvme_docker              latest            187c3d727248      11 days ago       28.7GB
stable-diffusion-inf      latest            de3091668a4c      11 days ago       21.1GB
```

- testing-gpu (GPU image)

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
milestone_9	latest	15b2c97f6fb1	2 weeks ago	21.3GB
nvme/api	latest	1371a8c94b44	2 weeks ago	28.7GB
testing-gpu	latest	6f9e22cc624b	3 weeks ago	7.15GB
milestone-7	latest	fad40770b96f	4 weeks ago	21.4GB
milestone-7-dev	latest	46f435f2b12f	4 weeks ago	21.4GB

3.3 Check Docker Network

- Containers need to share a common network to communicate (your GUI talks to the GPU helper over this network).
- We're using gui_network as that shared channel.

Ensure the gui_network exists by running:

```
docker network ls
```

This lists all Docker networks so you can see if gui_network is already there.

- If you don't see gui_network in the list, create it using:
 - You must create it so your containers can connect to each other.
 - Once it's created, any container you attach to gui_network will automatically be able to talk to the others.

```
docker network create gui_network
```

4. Run GUI Container

- It starts the full GUI application (front end, back end, database, and inference) all at once inside a single container.
- It gives the container access to your GPUs and fast NVMe storage so that heavy processing runs smoothly.
- It opens the right network ports so you can reach the interface in your browser—and if it ever crashes, Docker will restart it automatically.

Use the following command to start your GUI container:

```
sudo docker run -d \  
  --name <container_name> \           # Unique name for the container  
  --restart=unless-stopped \  
  --gpus "<specified_gpus>" \         # GPUs to expose (e.g. "device=0,1")  
  -it \  
  --pid=host \  
  --ipc=host \  
  --privileged=true \  
  --ulimit memlock=-1 \  
  --ulimit stack=67108864 \  
  --network gui_network \  
  -p <port_1>:5432 \                   # database port  
  -p <port_2>:3005 \                   # backend API port  
  -p <port_3>:80 \                     # frontend port  
  -p <port_4>:5004 \                   # inference port  
  -p <port_5>:8000 \                   # VLLM host port  
  -v /dev:/dev \  
  -v /models:/models \  
  -v <mount_point>:<mount_point> \     # NVMe mount path inside container  
  -v /media:/media \  
  -v /proc:/proc \  
  -v /sys:/sys \  
  -v /var/run/docker.sock:/var/run/docker.sock \  
  -v /etc/os-release:/etc/os-release:ro \  
  -v /usr/local/datasets:/usr/local/datasets \  
  -v /usr/local/models:/usr/local/models \  
  -v /dev/mapper:/dev/mapper \  
  -e PORT=<port_2> \  
  -e INFERENCE=<port_4> \  
  -e VLLM_HOST=<port_5> \  
  -e NVME_PATH=<mount_point> \  
  -e SPECIFIED_GPUS="<specified_gpus>" \  
  <image_name>                        # Docker image to launch
```

4.1 Parameters to Update:

- **<container_name>**
A unique name for your GUI container (e.g. miphi-gui, demo-1).
- **--gpus""device=0,1""**
Which GPU devices on your host you want to expose to the container ("device=0", "device=1", or "device=0,1", etc.).
- **-p <port_1>:5432**
Database port – maps your host DB port (e.g. 5432) into the contain
- **-p <port_2>:3005**
Backend API port – maps your host backend port (e.g. 3005) into the container.
- **-p <port_3>:80**
Frontend port – maps your host HTTP port (e.g. 80) into the container.
- **-p <port_4>:5004**
Inference port – maps your host inference-service port (e.g. 5004) into the container.
- **-p <port_5>:8000**
VLLM host port – maps your host VLLM service port (e.g. 8000) into the container.
- **-v <mount_point>:<mount_point>**
NVMe mount – e.g. if your ultra-fast storage is mounted at /mnt/nvme0 on the host, use -v /mnt/nvme0:/mnt/nvme0.
This gives your container direct access to that NVMe path.
- **-e NVME_PATH=<mount_point>**
Set this to your same NVMe path (e.g. NVME_PATH=/mnt/nvme0) so your app knows where to find it inside the container.
- **<image_name>**
The Docker image you loaded (e.g. integration-stream, testing-gpu, or your custom).

4.2 Example Parameter Values:

Replace the placeholders in your template with the values shown :

- <container_name> → integration-stream
- <image_name> → integration-stream
- "<specified_gpus>" → "device=0,1,2,3,4,5,6,7"
- <port_1> → 5794
- <port_2> → 3341
- <port_3> → 7068
- <port_4> → 5341
- <port_5> → 8089
- <mount_point> → /mnt/nvme1n1p1
- NVME_PATH=<mount_point> → NVME_PATH=/mnt/nvme1n1p1

```
miphi@miphi:~$ sudo docker run -d \
--name integration-stream \
--restart=unless-stopped \
--gpus '"device=0,1,2,3,4,5,6,7"' \
-it \
--pid=host \
--ipc=host \
--privileged=true \
--ulimit memlock=-1 \
--ulimit stack=67108864 \
--network gui_network \
-p 5794:5432 \
-p 3341:3005 \
-p 7068:80 \
-p 8089:8000 \
-p 5341:5004 \
-v /dev:/dev \
-v /models:/models \
-v /mnt/nvme1n1p1:/mnt/nvme1n1p1 \
-v /media:/media \
-v /proc:/proc \
-v /sys:/sys \
-v /var/run/docker.sock:/var/run/docker.sock \
-v /etc/os-release:/etc/os-release:ro \
-v /usr/local/datasets:/usr/local/datasets \
-v /usr/local/models:/usr/local/models \
-v /dev/mapper:/dev/mapper \
-e PORT=3341 \
-e INFERENCE=5341 \
-e VLLM_HOST=8089 \
-e NVME_PATH=/mnt/nvme1n1p1 \
-e SPECIFIED_GPUS="0,1,2,3,4,5,6,7" \
integration-stream
81e6155f8b0f0070af8746838999a79cd456f0af1501ef98b2e93fc872671cb2
miphi@miphi:~$
```

4.3 Verify Container Status

- Open your terminal and run

```
docker ps
```

- This command will display a table of all currently running containers, showing their names, IDs, and status.
- In the output, find the row with integration-stream under the NAMES column and check that its STATUS is “Up.” That confirms your GUI container is running correctly.

```
22 additional security updates can be applied with ESM Apps.
Learn more about enabling ESM Apps service at https://ubuntu.com/esm

Last login: Mon May 12 11:50:00 2025 from 192.168.100.73
miphi@miphi:~$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
81e6155f8b0f	integration-stream	"/docker-entrypoint..."	58 minutes ago	Up 58 minutes	0.0.0.0:7068->80/tcp, [::]:7068->80/tcp, 0.0.0.0:3341->3005/tcp, [::]:3341->3005/tcp, 0.0.0.0:5341->5004/tcp, [::]:5341->5004/tcp, 0.0.0.0:5794->5432/tcp, [::]:5794->5432/tcp, 0.0.0.0:8089->8000/tcp, [::]:8089->8000/tcp
4e6a406a1edb	ubuntu:22.04	"/bin/bash"	29 hours ago	Up 19 hours	
e3f92a808e5c	nvme_docker	"/docker-entrypoint..."	2 days ago	Up 2 days	0.0.0.0:7034->80/tcp, [::]:7034->80/tcp, 0.0.0.0:3204->3005/tcp, [::]:3204->3005/tcp, 0.0.0.0:5304->5009/tcp, [::]:5304->5009/tcp, 0.0.0.0:5602->5432/tcp, [::]:5602->5432/tcp

5. Run GPU Utility Container

- It launches a small helper container that has direct access to **all** your GPUs for easy monitoring and testing.
- It opens a host port (e.g. 9999) so you can connect to the GPU utility service from your machine or other tools.

```
sudo docker run -d \
  --gpus all \
  --network gui_network \
  -p <port_no>:9999 \           # host port for GPU utility service
  --name testing-gpu-container \ # container name for easy reference
  <image_gpu_name>
```

5.1 Parameters to Update:

- **-p<port_no>:9999**
Host port for the GPU utility service (e.g. 9999).
- **<image_gpu_name>**
The name of your GPU utility image (e.g. testing-gpu).

5.2 Example Parameter Values:

- <port_no> → 9999
- <image_gpu_name> → testing-gpu

```
miphi@miphi:~$ sudo docker run -d \  
--gpus all \  
--network gui_network \  
-p 9999:9999 \  
--name testing-gpu-container \  
testing-gpu  
[sudo] password for miphi:  
adbc1cd1d5baaa240764deccd8df088fdea681f58cd593a0bedd6adb0135caff  
miphi@miphi:~$
```

5.3 Verify GPU Container:

- Run **docker ps** and look for testing-gpu-container in the list.
- Check that its STATUS is “Up” to confirm the GPU helper is running.

```
Run 'do-release-upgrade' to upgrade to it.  
Last login: Mon May 12 16:48:23 2025 from 192.168.100.75  
antpc@antpc:~$ docker ps  
CONTAINER ID   IMAGE          COMMAND                  CREATED        STATUS        PORTS                               NAMES  
bd08ddd4bf54   testing-gpu    "/start.sh"             3 weeks ago   Up 4 days    0.0.0.0:9999->3000/tcp, [::]:9999->3000/tcp   testing-gpu-container
```

6. Post-Deployment Checklist

6.1 Verify running containers

- Run **docker ps** and ensure both your GUI and GPU containers are listed and in the “Up” state.

```
22 additional security updates can be applied with ESM Apps.
Learn more about enabling ESM Apps service at https://ubuntu.com/esm

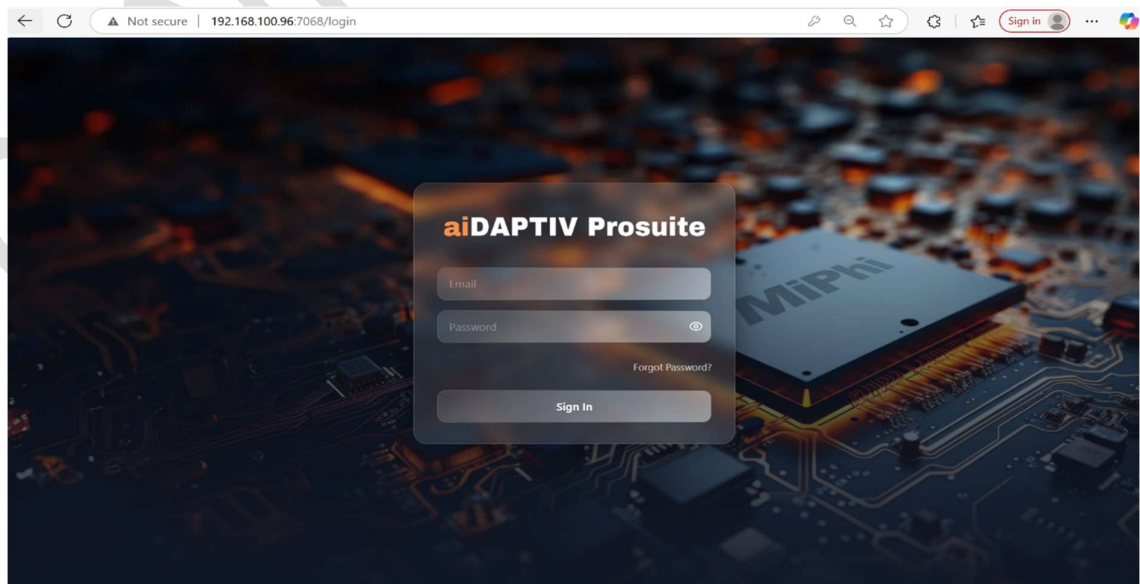
Last login: Mon May 12 11:50:00 2025 from 192.168.100.73
miphi@miphi:~$ docker ps
CONTAINER ID   IMAGE                                COMMAND                                CREATED        STATUS        PORTS
81e6155f8b0f   integration-stream                  "/docker-entrypoint..."            58 minutes ago Up 58 minutes 0.0.0.0:7068->80/tcp,
[::]:7068->80/tcp, 0.0.0.0:3341->3005/tcp, [::]:3341->3005/tcp, 0.0.0.0:5341->5004/tcp, [::]:5341->5004/tcp, 0.0.0.0:5794->54
32/tcp, [::]:5794->5432/tcp, 0.0.0.0:8089->8000/tcp, [::]:8089->8000/tcp   integration-stream
4e6a406a1edb   ubuntu:22.04                       "/bin/bash"                           29 hours ago   Up 19 hours
e3f92a808e5c   nvme_docker                        "/docker-entrypoint..."            2 days ago     Up 2 days     0.0.0.0:7034->80/tcp,
[::]:7034->80/tcp, 0.0.0.0:3204->3005/tcp, [::]:3204->3005/tcp, 0.0.0.0:5304->5009/tcp, [::]:5304->5009/tcp, 0.0.0.0:5602->54
32/tcp, [::]:5602->5432/tcp   nvme-docker-1
```

6.2 Access the GUI

- In your docker ps output’s PORTS column, find the entry that maps to 80/tcp (e.g. 0.0.0.0:7068->80/tcp); the number before :80 (here, 7068) is your front-end port

`http://<server-ip>:<gui_port>`

(e.g. <http://192.168.100.96:7068>)



6.3 Confirm GPU access

- Execute inside your GUI container:

```
docker exec -it <container_name> nvidia-smi
```

- You should see your GPU(s) listed with their status and memory usage.

```
miphi@miphi:~$ docker exec -it integration-stream nvidia-smi
Mon May 12 18:56:14 2025
```

NVIDIA-SMI 550.120				Driver Version: 550.120				CUDA Version: 12.4			
GPU	Name	Perf	Persistence-M	Bus-Id	Disp.A	Memory-Usage	Volatile	Uncorr.	ECC		
Fan	Temp		Pwr:Usage/Cap				GPU-Util	Compute	M. MIG M.		
0	NVIDIA RTX A6000	P8	Off	00000000:01:00.0	Off	4175MiB / 49140MiB	0%	Default	Off N/A		
30%	38C		20W / 300W								
1	NVIDIA RTX A6000	P8	Off	00000000:21:00.0	Off	1639MiB / 49140MiB	0%	Default	Off N/A		
31%	41C		19W / 300W								
2	NVIDIA RTX A6000	P2	Off	00000000:41:00.0	Off	9799MiB / 49140MiB	45%	Default	Off N/A		
31%	58C		134W / 300W								
3	NVIDIA RTX A6000	P8	Off	00000000:61:00.0	Off	4MiB / 49140MiB	0%	Default	Off N/A		
31%	44C		23W / 300W								
4	NVIDIA RTX A6000	P8	Off	00000000:81:00.0	Off	4MiB / 49140MiB	0%	Default	Off N/A		
30%	43C		21W / 300W								
5	NVIDIA RTX A6000	P8	Off	00000000:A1:00.0	Off	4MiB / 49140MiB	0%	Default	Off N/A		
31%	43C		22W / 300W								
6	NVIDIA RTX A6000	P8	Off	00000000:C1:00.0	Off	4MiB / 49140MiB	0%	Default	Off N/A		
30%	43C		19W / 300W								
7	NVIDIA RTX A6000	P8	Off	00000000:E1:00.0	Off	10837MiB / 49140MiB	0%	Default	Off N/A		
30%	39C		23W / 300W								

Processes:							GPU Memory Usage
GPU	GI ID	CI ID	PID	Type	Process name		
0	N/A	N/A	2420631	C	python3	4102MiB	
1	N/A	N/A	1313709	C	python3	1632MiB	
2	N/A	N/A	2497063	C	/opt/venv/bin/python3	9772MiB	
7	N/A	N/A	3502297	C	python3	10766MiB	